

AI AGENT FUNDAMENTALS

Six-Project Practical Exercise — Summary Report

Project 1: The Worker / Tool-Location Test

Goal: Understand the difference between where the AI “thinks” (agent-loop location) and where a tool actually runs (tool-execution location).

What Was Done:

- Started a task in the cloud using only a web-reachable source, then closed and reopened it from another device — the task remained complete.
- Connected a local folder on the desktop app and asked Claude to read a file from it.
- Tested access after moving the file and after disconnecting the folder connector, to see how local access behaves.

Result / Finding:

- Cloud-only tasks do not depend on any single device; the reasoning happens on the server, so closing a tab does not stop the work.
- Local-folder tasks depend on live, real-time access to that folder — if the file moves or access is removed, the task fails or reports the file missing.
- Part 1: <https://claude.ai/share/483b3ae5-dca1-4384-83f2-79e1197cecae> (Chat)
- Part 2: Screenshot has been attached for proof.

Status: **Complete**

Project 2: The Three-Tier Audit

Goal: Trace a real task from where its context lives to where its final output is stored, and identify what would be lost if the account disappeared.

What Was Done:

- Asked Claude to summarize a website into a PDF, in a normal chat session.
- Saved the finished PDF to Google Drive using the Google Drive connector.

Result / Finding:

- State layer used: Session (a plain chat, no saved Project or memory involved).
- Final file tier: Tier 3 — a real file saved in Google Drive, independent of the vendor account.
- If the account were deleted, the chat history would be lost, but the PDF in Google Drive would remain safe.
- Claude Chat URL: <https://claude.ai/share/20c7bdfd-cd4d-45d2-b328-ba3a8913985b>

Status: **Complete**

Project 3: The Gate Lab

Goal: Observe the difference between actions an AI can take automatically and actions that require explicit approval.

What Was Done

- Asked Claude to draft (not send) a meeting-reminder email through the Gmail connector.
- Asked Claude to then send that draft, to see if approval was requested.

Result / Finding

- Allowed automatically: creating the email draft — no approval was requested.
- Escalated for approval: sending the email — Claude asked for confirmation through Gmail before sending.
- Chosen approach for real workflows: let the AI draft freely, but always review and approve before anything is actually sent, since sending is irreversible.
- Claude Chat URL: <https://claude.ai/share/0ac533de-2bdd-429a-9292-b19733c94f83>

Status: **Complete**

Project 4: The First Cloud Schedule

Goal: Build a recurring cloud task that runs on its own schedule, independent of any device being on.

Six Answers Defined Before Scheduling

- Trigger: A fixed time each day / a fixed day and time each week.
- Touch: Web search plus the Gmail and Google Drive connectors — all cloud-reachable, no local folder used.
- Device independence: The task must run whether or not the computer is on.
- Success signal: A PDF appears in Google Drive (daily) and an email arrives (weekly).
- Autonomy: Runs fully automatically, with approvals set to always allow for this low-risk task.
- Empty case: If there is no major news, the task still produces a short PDF saying so, instead of skipping the run.

What Was Done

- Built a scheduled task in Claude Cowork: daily research and PDF saved to Google Drive, plus a Wednesday email with a weekly summary.
- Tested the workflow manually before scheduling, then left the computer switched off at the scheduled run time on two separate days.

Result / Finding

- The task fired successfully on its own, with the computer off, on two separate occasions.
- Each run left a verifiable Tier 3 output — a PDF saved to Google Drive.
- Screenshot has been attached as a proof.

Status: **Complete**

Project 5: One Task, Two Harnesses

Goal: Give the identical task to two different AI products and compare how each implements the same underlying structure.\

What Was Done

The same prompt was given to both ChatGPT Work and Claude Cowork: search the web for today's top 3 AI industry news stories, summarize each, save as a PDF, show reasoning, and ask for approval before finalizing.

Comparison

Concept	ChatGPT Work	Claude Cowork
Heartbeat	Started immediately after the prompt was sent.	Started immediately after the prompt was sent (same).
Reach	Ran a smaller number of broader web searches.	Ran multiple targeted searches per company, plus page fetches for verification.
Loop	Reasoning steps were mostly hidden; went straight to a request for approval.	Reasoning was shown step by step, including why each story was chosen.
Gate	Asked for approval, then built and delivered the PDF right away.	Asked for approval and waited; only built the PDF after an explicit "yes."
Body	Delivered the PDF directly after approval.	Built the PDF using a document skill, visually verified it, then delivered it.

Conclusion: Both products followed the same six-part shape (heartbeat, reach, loop, spine, gate, body), but implemented at least three of those parts differently — most notably how strictly each one honored the approval gate before producing the file.

ChatGPT Chat URL: <https://chatgpt.com/share/6aa07de4-380c-83e9-8beb-78ec35d037c6>

Claude Screenshot has been attached:

Status: **Complete**

Project 6: The Portability Drill

Goal: Assume the AI vendor disappears tomorrow. Using only Tier 3 files and personally-owned notes, work out what of a real workflow could still be rebuilt — and what could not.

Workflow Used For The Drill

The Project 4 daily/weekly tech-news workflow: its instructions, its cloud sources, its definition of done, and its last saved PDF.

Reconstruction Check

Could Be Reconstructed (Tier 3 / Own Notes)	Could NOT Be Reconstructed (Lock-in Risk)
Original prompt text, once saved outside the chat.	Exact scheduled-task configuration (frequency, approval mode setting).
Final PDF outputs already saved to Google Drive.	The exact search-and-reasoning process used to research and choose stories.
Definition of done, once written down (e.g. “PDF appears in Drive”).	Connector permissions (Gmail, Google Drive) — would need to be reconnected from scratch on another tool.
General source description (“web search for tech news”).	Full conversation history and context built up across the chat.

Lock-in Exposure (Next Portability Backlog)

- Write down the exact scheduled-task settings (times, approval mode) outside of Cowork, not just in the product itself.
- Keep a short written definition of done for each recurring task, separate from the chat.
- Expect to fully reconnect any connectors (Gmail, Drive) from scratch on a different AI product.

Status: **Complete**

Overall Summary:

Project	Focus	Status
1	Agent-loop location vs. tool-execution location	Complete
2	Tracing state layers and file tiers	Complete
3	Automatic actions vs. approval gates	Complete
4	Building a reliable cloud schedule	Complete
5	Comparing two AI products on the same task	Complete
6	Planning for vendor independence	Complete

Sharmeen Fatima
Founder & CEO
CreativeDev
www.creativedev.store